

PThreads – An Introduction

CS F422 Parallel Computing Tutorial

PThreads

- POSIX Thread
- POSIX – A standard for Unix like OS
- Library linked with C programs

Things to remember

- Threads within a process share same address space
 - Data, open files, current working directory ...
- Each thread has a unique
 - Thread id, registers, stack ...

Execution

- `#include <pthread.h>`
- `gcc -g -Wall -o outputfilename filename.c -lpthread`

E.g. 1

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *print_message_function( void *ptr );

main() {
    pthread_t thread1, thread2;
    char *message1 = "Thread 1";
    char *message2 = "Thread 2";
    int  iret1, iret2;

    /* Create independent threads each of which will execute function */

    iret1 = pthread_create( &thread1, NULL, print_message_function, (void*) message1);
    iret2 = pthread_create( &thread2, NULL, print_message_function, (void*) message2);

    /* Wait till threads are complete before main continues. Unless we
    /* wait we run the risk of executing an exit which will terminate
    /* the process and all threads before the threads have completed. */

    pthread_join( thread1, NULL);
    pthread_join( thread2, NULL);

    printf("Thread 1 returns: %d\n",iret1);
    printf("Thread 2 returns: %d\n",iret2);
    exit(0);
}

void *print_message_function( void *ptr ){
    char *message;
    message = (char *) ptr;
    printf("%s \n", message);
}
```

Starting Threads

- `int pthread_create(
 pthread_t *thread ,
 const pthread_attr_t *attr,
 void *(*start_routine)(void*),
 void *arg_p);`
 - `thread` – return thread id
 - `attr` – set to `NULL`. Attributes include detached state, scheduling policy, scheduling parameter, inheritsched attribute, scope ...
 - `Void *(*start_routine)` – pointer to function to be threaded with single pointer to void argument
 - `*arg` – pointer to argument of function

Stopping Threads

- `int pthread_join(
pthread_t thread,
void** ret_val_p);`
- Output E.g. 1
Thread 1
Thread 2
Thread 1 returns: 0
Thread 2 returns: 0

Killing threads

- `void pthread_exit(void *retval)`
- Kills the thread
- Return value of thread stored in `retval`, which is not of local scope

Synchronization

- Mutexes
- Joins
- Condition variables

E.g. 2 Mutexes

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *functionC();
pthread_mutex_t mutex1 = PTHREAD_MUTEX_INITIALIZER;
int counter = 0;

main(){
    int rc1, rc2;
    pthread_t thread1, thread2;

    /* Create independent threads each of which will execute functionC */

    if( (rc1=pthread_create( &thread1, NULL, &functionC, NULL)) )
        printf("Thread creation failed: %d\n", rc1);

    if( (rc2=pthread_create( &thread2, NULL, &functionC, NULL)) )
        printf("Thread creation failed: %d\n", rc2);

    /* Wait till threads are complete before main continues. Unless we */
    /* wait we run the risk of executing an exit which will terminate */
    /* the process and all threads before the threads have completed. */

    pthread_join( thread1, NULL);
    pthread_join( thread2, NULL);

    exit(0);
}

void *functionC(){
    pthread_mutex_lock( &mutex1 );
    counter++;
    printf("Counter value: %d\n",counter);
    pthread_mutex_unlock( &mutex1 );
}
```

Mutexes

- Prevent data inconsistencies due to race conditions (same memory changed by threads, but output depends on order of operations)
- Serialize shared resources
- In E.g. 2, if locking was not applied, the two threads may have written same result as 1, but with locking correct result is obtained, i.e., 2.
- Output E.g. 2
Counter value: 1
Counter value: 2

E.g. 3 Joins

```
#include <stdio.h>
#include <pthread.h>

#define NTHREADS 10
void *thread_function(void *);
pthread_mutex_t mutex1 = PTHREAD_MUTEX_INITIALIZER;
int counter = 0;

main(){
    pthread_t thread_id[NTHREADS];
    int i, j;

    for(i=0; i < NTHREADS; i++)
        pthread_create( &thread_id[i], NULL, thread_function, NULL );

    for(j=0; j < NTHREADS; j++)
        pthread_join( thread_id[j], NULL);

    /* Now that all threads are complete I can print the final result.      */
    /* Without the join I could be printing a value before all the threads */
    /* have been completed.                                                */

    printf("Final counter value: %d\n", counter);
}

void *thread_function(void *dummyPtr){
    printf("Thread number %ld\n", pthread_self());
    pthread_mutex_lock( &mutex1 );
    counter++;
    pthread_mutex_unlock( &mutex1 );
}
```

Joins

- Wait for a thread to finish
- Say master thread waiting for all threads to complete and then terminate
- In E.g. 3, main thread waits for all to complete and then prints the counter value
- Output E.g. 3
 - Thread number 1026
 - Thread number 2051 ...
 - Final counter value: 10

E.g. 4

Condition Variables

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

pthread_mutex_t count_mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_mutex_t condition_mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t condition_cond = PTHREAD_COND_INITIALIZER;

void *functionCount1();
void *functionCount2();
int count = 0;
#define COUNT_DONE 10
#define COUNT_HALT1 3
#define COUNT_HALT2 6

main(){
    pthread_t thread1, thread2;

    pthread_create( &thread1, NULL, &functionCount1, NULL);
    pthread_create( &thread2, NULL, &functionCount2, NULL);
    pthread_join( thread1, NULL);
    pthread_join( thread2, NULL);

    exit(0);
}

void *functionCount1(){
    for(;;){
        pthread_mutex_lock( &condition_mutex );
        while( count >= COUNT_HALT1 && count <= COUNT_HALT2 )
            pthread_cond_wait( &condition_cond, &condition_mutex );
        pthread_mutex_unlock( &condition_mutex );

        pthread_mutex_lock( &count_mutex );
        count++;
        printf("Counter value functionCount1: %d\n",count);
        pthread_mutex_unlock( &count_mutex );
    }
}

void *functionCount2(){
    for(;;){
        pthread_mutex_lock( &condition_mutex );
        if( count < COUNT_HALT1 || count > COUNT_HALT2 )
            pthread_cond_signal( &condition_cond );
        pthread_mutex_unlock( &condition_mutex );

        pthread_mutex_lock( &count_mutex );
        count++;
        printf("Counter value functionCount2: %d\n",count);
        pthread_mutex_unlock( &count_mutex );

        if(count >= COUNT_DONE) return(NULL);
    }
}
```

Condition Variables

- Variable of type `pthread_cond_t`
- Used with appropriate functions for waiting and later, process continuation
- Condition variable associated with a mutex to avoid race condition
- Functions used in conjunction with condition variable
 - Creating/destroying
 - `pthread_cond_init`
 - `pthread_cond_t cond = PTHREAD_COND_INITIALIZER;`
 - `pthread_cond_destroy`

Condition Variables

- Waiting on condition
 - `pthread_cond_wait`
 - `pthread_cond_timedwait` – limit on blocking time
- Waking threads based on condition
 - `pthread_cond_signal`
 - `pthread_cond_broadcast` – wake up all threads blocked by the specified condition variable
- In E.g. 4, `functionCount1()` halted while count was between `COUNT_HALT1` and `COUNT_HALT2`. `functionCount2()` increments count during this time. Everything else is random

Condition Variables

- Output E.g. 4

Counter value functionCount1: 1

Counter value functionCount1: 2

Counter value functionCount1: 3

Counter value functionCount2: 4

Counter value functionCount2: 5

Counter value functionCount2: 6

Counter value functionCount2: 7

Counter value functionCount1: 8

Counter value functionCount1: 9

Counter value functionCount1: 10

Counter value functionCount2: 11

Reference

- [1] An Introduction to Parallel Programming by Peter S. Pacheco
- [2] [POSIX thread \(pthread\) libraries](#)